

Collection Concurrency



Robert Horvick

SOFTWARE ENGINEER

@bubbafat www.roberthorvick.com



Overview



Concurrency

- Overview
- Problems

Solutions

- Caller synchronization
- Monitor locking
- Read/write locking

.NET concurrent collections



Concurrency

Two or more operations executing at the same time (concurrently).



Concurrency



Multiple threads executing within a single process accessing a shared resource (e.g., a shared `List<T>`)



Multiple processes running on the same computer accessing a shared resource (e.g., a shared file)



Multiple processes running on different computers accessing a shared resource (e.g., a shared database table)



```
class Job {  
    readonly int Priority; // set in constructor  
    public void Process() {  
        ...  
    }  
}
```

Job Class

A basic job class that contains a priority and a process method.



```
var jobs = new PriorityQueue<Job>();
```

```
for(int i = 0; i < 100; i++) {  
    jobs.Enqueue(new Job(i));  
}
```

```
while(jobs.Count > 0) {  
    var job = jobs.Dequeue();  
    job.Process();  
}
```

◀ The priority queue of jobs to execute

◀ Add 100 jobs to the queue (the constructor parameter is the job priority)

◀ Dequeue and process each of the job in priority order.



```
Thread[] adders = new Thread[4];
```

```
ThreadStart addJobs = delegate() {  
    for(int i = 0; i < 25; i++) {  
        jobs.Enqueue(new Job(i));  
    }  
}
```

```
for(int i=0; i < adders.Length; i++) {  
    adders[i] = new Thread(addJobs);  
    adders[i].Start();  
}
```


◀ Add jobs using 4 threads

◀ Each thread will add 25 jobs to the queue


◀ Create the 4 threads and start them to add the jobs to the queue concurrently



Add new value to end of
heap




Add new value to end of
heap




Add new value to end of heap

While the heap property is not satisfied, swap with its parent




Add new value to end of heap

While the heap property is not satisfied, swap with its parent




Add new value to end of heap

While the heap property is not satisfied, swap with its parent



Thread 1


Thread 2



Thread 1

Thread 2


7




Thread 1



Thread 2

7



Thread 1	Thread 2
7	25






Thread 1



Thread 2


25



Thread 1

Thread 2


7



Thread 1

Thread 2


7




Thread 1

Thread 2

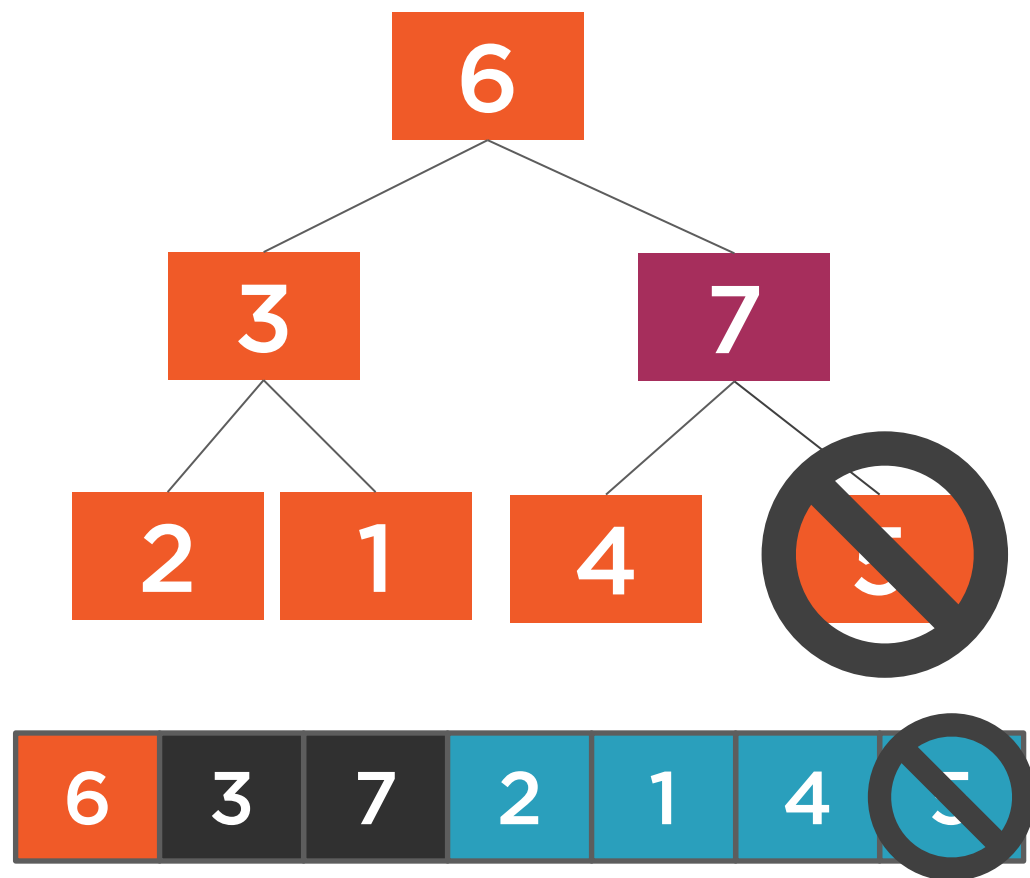
7





Thread 1	Thread 2
7	25



Thread 1	Thread 2
7	25







Concurrent updates to non-
concurrency-safe
collections can lead to
unexpected behavior and
data loss



Thread 1

Thread 2





Dequeue



Thread 1

Thread 2

100





Dequeue



Thread 1

Thread 2

100





Dequeue



Thread 1

Thread 2

100





Dequeue



Thread 1

Thread 2

100



Dequeue




Thread 1

100

Thread 2

4



Dequeue



Caller Synchronization


The caller is responsible for ensuring all access to the collection is performed in a concurrency-safe manner




Caller Synchronization



Allows concurrent-safe access to non-concurrency-safe collections



No overhead when the collection is used non-concurrently



The caller can determine the optimal synchronization approach



```
object jobsLock = new object();
```

```
// ...
```

```
lock(jobsLock) {
```

```
    jobs.Enqueue(new Job(...));
```

```
}
```

```
lock(jobsLock) {
```

```
    Job nextJob = jobs.Dequeue();
```

```
    nextJob.Process();
```

```
}
```

◀ Create the object that will be used as the shared monitor lock object

◀ Take the lock before enqueueing jobs

◀ Lock before dequeuing the next job

◀ The job can be processed



```
object jobsLock = new object();  
  
// ...  
  
while(jobs.Count > 0) {  
  
    lock(jobsLock) {  
        Job nextJob = jobs.Dequeue()  
        nextJob.Process();  
    }  
}
```

◀ **Count needs to be called within the same lock scope as the call to Dequeue**

◀ **The Process method is holding the lock open.**



```
while(jobs.Count > 0) {  
    Job nextJob = null;  
  
    lock(jobsLock) {  
        if(jobs.Count > 0) {  
            nextJob = jobs.Dequeue();  
        }  
    }  
  
    if(nextJob != null) {  
        nextJob.Process();  
    }  
}
```

- ◀ Check if the count is more than 0 while outside the lock
- ◀ Take the lock
- ◀ Check the count again while under the lock
- ◀ While holding the lock, Dequeue the next job
- ◀ If we dequeued a job, then process it



Caller Synchronization Using Monitor Lock

Pros

Non-concurrent-safe collections can be used in concurrent environments

Easy to implement

Cons

Caller is responsible for all thread synchronization

Readers block other readers

Easy to implement wrong



**CODE – Caller Non-Locking
Client and then locking client**



Collection Synchronization

The caller is responsible for ensuring all access to the collection is performed in a concurrency-safe manner



Collection Synchronization

Monitor Locking

A single monitor lock is used to serialize access to the container

Monitor locks are very light-weight

Readers block other readers

Reader/Writer Locking

A single reader/writer lock is used to serialize write-access to the container

The reader/writer lock allows multiple readers concurrently while blocking writes

Concurrent reading can overcome performance costs versus monitors



Collection Synchronization

Monitor Locking

A single monitor lock is used to serialize access to the container

Monitor locks are very light-weight

Readers block other readers

Reader/Writer Locking

A single reader/writer lock is used to serialize write-access to the container

The reader/writer lock allows multiple readers concurrently while blocking writes

Concurrent reading can overcome performance costs versus monitors

The caller must still avoid non-concurrency-safe access patterns



```
object syncLock = new object();  
// ...  
  
public void Enqueue(T value) {  
    lock(syncLock) {  
        heap.Push(value);  
    }  
}
```

- ◀ A single synchronization object is used to serialize access to the priority queue
- ◀ The lock is taken during any operation that requires access to the heap



CODE – Locking Queue



Reader Writer Locks

The .NET ReaderWriterLockSlim class used to provide concurrent readers while serializing all writers.



```
var rwLock = new ReaderWriterLockSlim();  
  
// ...  
  
public void Enqueue(T value) {  
    rwLock.EnterWriteLock();  
  
    try  
    {  
        heap.Push(value);  
    }  
  
    finally  
    {  
        rwLock.ExitWriteLock();  
    }  
}
```

- ◀ A single `ReaderWriterLockSlim` instance serializes writes and blocks writes while allowing concurrent reads.
- ◀ The write lock is entered before a non-concurrency-safe operation. All reads and writes are blocked until this is exited.
- ◀ The non-concurrent-safe operation runs within a try-block
- ◀ In the finally-block the write lock is exited



```
var rwLock = new ReaderWriterLockSlim();

// ...

public T Peek() {
    rwLock.EnterReadLock();

    try
    {
        return heap.Top();
    }

    finally
    {
        rwLock.ExitReadLock();
    }
}
```

- ◀ In a read-only method a read lock is used to allow concurrent readers while blocking writes
- ◀ The read operation is performed within a try-block
- ◀ When the number of readers is zero then writes will be allowed again



CODE – RW Locking Queue



Concurrent .NET Collections

ConcurrentDictionary<TK,TV>

ConcurrentQueue<T>

ConcurrentStack<T>

ConcurrentBag<T>



.NET Concurrent Collections



Not drop-in replacements for existing collection types



Prefer these types with code requiring concurrency-safe collections



`ConcurrentQueue` and `ConcurrentStack` are lock-free collections



```
using System.Collections.Concurrent;

//...

var queue = new ConcurrentQueue<int>();

queue.Enqueue(1);

int value;

if(queue.TryPeek(out value)) {
    Console.WriteLine(value);
}

if(queue.TryDequeue(out value)) {
    Console.WriteLine(value);
}
```

- ◀ The concurrent collections are in the `System.Collections.Concurrent` namespace
- ◀ Allocate a concurrent queue of integers
- ◀ Enqueue works the same as `Queue<T>`
- ◀ Peeking requires the “Try” pattern which avoids having to fail if the queue is empty
- ◀ Dequeueing also uses the “Try” pattern to avoid failure when the queue is empty

